

# Question 1: Understanding Hardware Configuration via /proc Filesystem

## *Command:*

```
more /proc/cpuinfo
```

This command provides detailed information about your CPU. Key terms to note:

- **processor:** Represents the logical processors (threads) in your system.
- **cores:** Refers to the physical cores of the CPU. Each core may support multiple threads if hyperthreading is enabled.

## *Verify with Lscpu:*

```
lscpu
```

- `lscpu` displays a concise summary of CPU architecture:
  - Core count
  - Thread count
  - Hyperthreading status
- **Hyperthreading:** Allows a single physical core to function as multiple logical processors.

## *Tips:*

- Use `lscpu` to answer subquestions (b) to (e) regarding cores, processors, frequency, and architecture.
- For parts (f) to (h), use these additional commands:
  - Total and free memory: `free` or `cat /proc/meminfo`
  - Forks and context switches: `grep` specific fields in `/proc/stat`.

## Question 2: Monitoring a Running Process Using top

### Steps:

#### 1. Compile and Run the Program:

```
gcc cpu.c -o cpu
./cpu
```

a. This runs the program cpu in an infinite loop.

#### 2. Open top in Another Terminal:

```
top
```

The top command displays a list of all running processes with their resource usage.

### Analyze the Output:

- **(a) Find the PID:**
  - Look for the process named cpu in the COMMAND column. The PID is shown in the first column.
- **(b) Check CPU and Memory Usage:**
  - %CPU: Percentage of CPU being used by the process.
  - %MEM: Percentage of memory consumed by the process.
- **(c) Observe the Process State:**
  - The STATE column indicates the current process state:
    - R: Running
    - S: Sleeping (Blocked)
    - Z: Zombie

### Exit Commands:

- Quit top: Press q.
- Stop the cpu program: Press Ctrl+C in the terminal where it is running.

## Question 3:

# Assignment Help: Shell Behavior - Child Processes and I/O Redirection

## 1. Compile and Run the Program

- **Task:** Compile and run the `cpu-print.c` program.
  - Use `gcc` to compile the program: `gcc cpu-print.c -o cpu-print`
  - Execute the compiled program: `./cpu-print`

The program will run indefinitely, printing output to the terminal.

## 2. Process Identification and Parent Processes

### *(a) Identify the PID of the `cpu-print` Process*

- Use the `ps` command to check the list of running processes:
  - `ps aux | grep <process-name>`
  - Look for the PID (Process ID) of `cpu-print` from the output.

### *(b) Trace Parent and Ancestor Processes*

- To find the parent process (PPID) of `cpu-print`:
  - `ps -o ppid= -p <PID-of-cpu-print>`
- Use `pstree` to trace the ancestors of the process:
  - `pstree -p <PPID>`

### 3. I/O Redirection to a File

#### *(c) Redirect the Output of cpu-print to a File*

- Run the program and redirect output to a file: `./cpu-print > /tmp/tmp.txt &`
- To explore file descriptors, use the `lsof` command: `lsof -p <PID-of-cpu-print>`
  - Look for entries for standard input (0), standard output (1), and standard error (2).
  - The output file (`/tmp/tmp.txt`) should be associated with 1 (stdout).

### 4. Pipe Implementation

#### *(d) Implement a Pipe and Examine Processes*

- Run the program with a pipe: `./cpu-print | grep hello &`
- Use `ps` to check the processes: `ps aux | grep -E "cpu-print|grep"`
- Check the file descriptors for each process: `lsof -p <PID-of-cpu-print>`  
`lsof -p <PID-of-grep>`

### 5. Investigate Built-in Commands

#### *(e) Explore cd, ls, history, and ps*

- To identify whether the commands are built-in or external executables: `which cd ls history ps`
- To verify if a command is a shell built-in or external: `type cd ls history ps`

## Important Notes:

- **Avoid Overwriting Files:** Be cautious when redirecting output to files. If running a command like `cpu-print` indefinitely, the file can grow large and fill up the disk.
- **Stop Running Processes:** Use `kill` to stop a process if needed: `kill -9 <PID>`
- **Check Disk Space:** To see how much disk space is available: `df -h`

## Question 4 : Virtual vs. Physical Memory Usage

### Overview

This question focuses on understanding how virtual and physical memory usage differ based on program behavior. The programs `memory1.c` and `memory2.c` allocate large arrays of the same size. However:

- **memory1.c** allocates the array but does not access it.
- **memory2.c** allocates the array and accesses/modifies its elements.

You will observe memory usage differences using the `ps` command and analyze why the **Resident Set Size (RSS)** varies when the array is accessed in `memory2.c`.

### Compiling and Running the Programs

1. Compile the programs using the following commands:

```
gcc memory1.c -o memory1
gcc memory2.c -o memory2
```

2. Run each program in separate terminals:

```
./memory1
```

```
./memory2
```

3. Each program will display its **Process ID (PID)**. Note down the PID for both programs.

## Observing Memory Usage

1. Open another terminal and use the `ps` command to check the memory usage of each program:

```
ps -p <PID> -o pid,vsz,rss,comm
```

Replace <PID> with the PID of `memory1` or `memory2`.

- a. **VSZ (Virtual Memory Size):** The total memory reserved for the program, including memory that may not yet be physically allocated.
- b. **RSS (Resident Set Size):** The amount of physical memory (RAM) currently allocated to the process.

## Q5: Help Guide for Disk Access Programs

### 1. Setting Up the Environment

- **Create Folder:** Create a directory named `disk-files`:

```
mkdir disk-files
```

- **Add Sample File:** Place foo.pdf inside the disk-files folder:

```
cp foo.pdf disk-files/
```

- **Generate Multiple Files:** Use the make-copies.sh script to create 5000 copies of foo.pdf with unique filenames in the folder. Make the script executable with the following command:

```
chmod +x make-copies.sh  
./make-copies.sh
```

## 2. Compiling and Running the Programs

- **Compilation:** Use GCC to compile disk.c and disk1.c:

```
gcc disk.c -o disk
```

- **Execution:** Run the program in the terminal:

```
./disk
```

- **Monitor Disk Usage:** Use tools like iostat, dstat, or iotop to monitor disk metrics.

## 3. Monitoring Disk Utilization

- Install iostat using your package manager.
- Use the following command to observe real-time disk activity: `iostat -d 1`

## 4. Clearing Disk Buffer Cache

To ensure you're reading from the disk and not from memory, clear the disk buffer cache using the following commands (Linux):

- **Clear page cache:**

```
sudo sync; sudo echo 1 > /proc/sys/vm/drop_caches
```

- **Clear dentries and inodes:**

```
sudo sync; sudo echo 2 > /proc/sys/vm/drop_caches
```

- **Clear both page cache and dentries/inodes:**

```
sudo sync; sudo echo 3 > /proc/sys/vm/drop_caches
```

## ***Question 6:***

### ***Debugging Programs Using GDB***

#### **Part 1: Debugging pointers.cpp**

1. **Compile with Debug Symbols:**



```
g++ -g pointers.cpp -o pointers
```

2. **Run GDB:**

```
gdb ./pointers
```

3. **Set a Breakpoint:**

```
break main
```

4. **Start the Program:**

```
run
```

5. **Step Through the Code:** Use next to execute line by line.

6. **Find the Faulty Line:** If the program crashes, use backtrace to locate the issue.

## Part 2: Debugging fibonacci.cpp

1. **Compile with Debug Symbols:**

```
g++ -g fibonacci.cpp -o fibonacci
```

2. **Run GDB:**

```
gdb ./fibonacci
```

3. **Set Breakpoints:**

a. At main: break main

b. Inside the loop (e.g., line 13): break <line\_number>

4. **Start the Program:**

```
run
```

5. **Print Variables:** Use `print <variable>` to observe variable values during execution.
6. **Step Through Code:** Use `next` to debug the loop and monitor logic.

## Common Commands Recap

- `break <line>`: Set a breakpoint.
- `run`: Start execution.
- `next`: Execute the next line.
- `print <variable>`: Show variable values.
- `backtrace`: Trace crash points.

## Question 7 : System Calls and strace

`strace` is a diagnostic, debugging, and instructional tool in Linux that allows users to trace the system calls made by a program during its execution. The objective is to trace the system calls made by an executable (e.g., a program or a Linux command like `ls`) using `strace`, analyze the trace output, and answer specific questions.

For example, to trace the system calls made by a program named `cpu.c`:

```
gcc cpu.c -o cpu.exe
```

```
strace <executable>
```

The output will contain detailed information about system calls made during the execution of the program or command.

```
<syscall_name>(<arguments>) = <return_value>
```

After running `strace`, count the number of lines representing system calls.

Find a list of supported system calls for your operating system. A good resource is the **man page for syscall** :

`man 2 <syscall_name>`